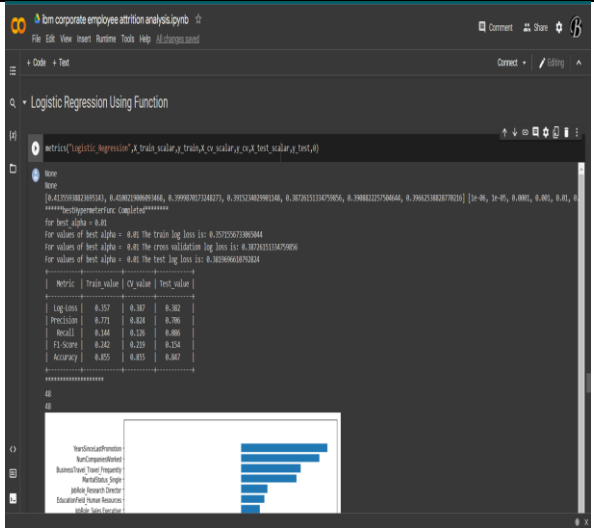
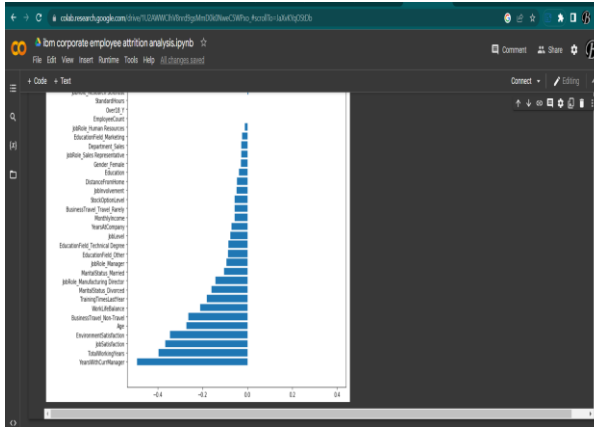


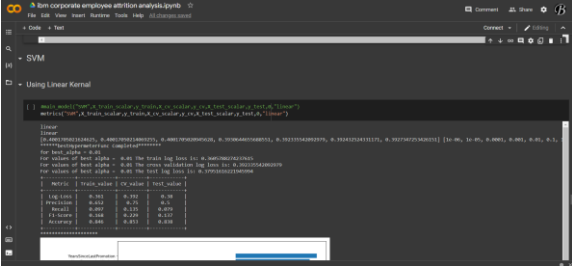
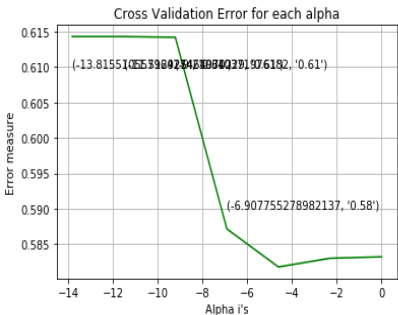
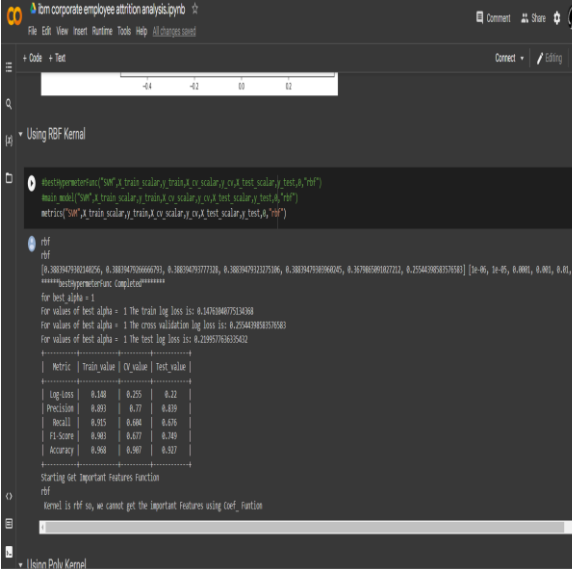
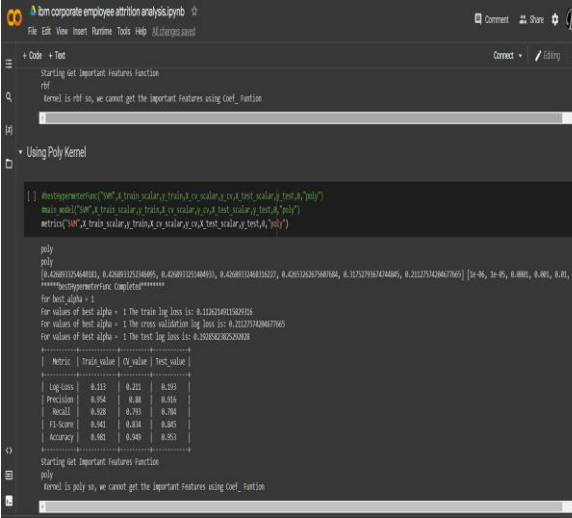
Project Development Phase Model Performance Test

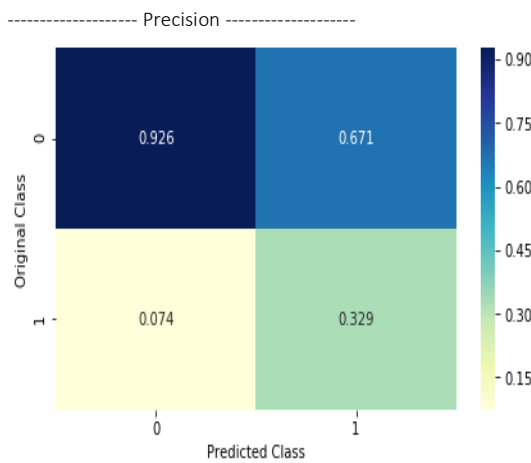
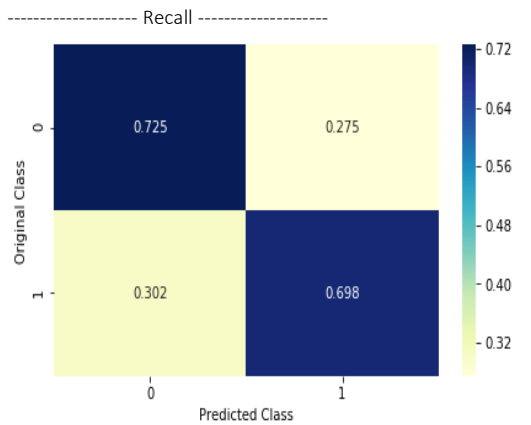
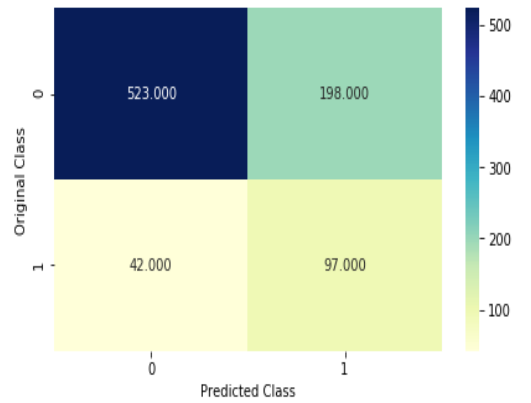
Date	10 November 2022
Team ID	PNT2022TMID37447
Project Name	Project - Corporate Employee Attrition Analytics
Maximum Marks	10 Marks

Model Performance Testing:

Project team shall fill the following information in model performance testing template.

S.No.	Parameter	Values	Screenshot																								
1.	Model Summary Regression analysis	<p>for best_alpha = 0.01 For values of best alpha = 0.01 The train log loss is: 0.3571556733065044 For values of best alpha = 0.01 The cross validation log loss is: 0.38726151334759856 For values of best alpha = 0.01 The test log loss is: 0.3819696610792824</p> <table border="1"> <thead> <tr> <th>Metric</th><th>Train_value</th><th>CV_value</th><th>Test_value</th></tr> </thead> <tbody> <tr> <td>Log-Loss</td><td>0.357</td><td>0.387</td><td>0.382</td></tr> <tr> <td>Precision</td><td>0.771</td><td>0.824</td><td>0.706</td></tr> <tr> <td>Recall</td><td>0.144</td><td>0.126</td><td>0.086</td></tr> <tr> <td>F1-Score</td><td>0.242</td><td>0.219</td><td>0.154</td></tr> <tr> <td>Accuracy</td><td>0.855</td><td>0.855</td><td>0.847</td></tr> </tbody> </table>	Metric	Train_value	CV_value	Test_value	Log-Loss	0.357	0.387	0.382	Precision	0.771	0.824	0.706	Recall	0.144	0.126	0.086	F1-Score	0.242	0.219	0.154	Accuracy	0.855	0.855	0.847	
Metric	Train_value	CV_value	Test_value																								
Log-Loss	0.357	0.387	0.382																								
Precision	0.771	0.824	0.706																								
Recall	0.144	0.126	0.086																								
F1-Score	0.242	0.219	0.154																								
Accuracy	0.855	0.855	0.847																								
	SVM Model	<p>linear linear [0.4001705021624625, 0.40017050214069255, 0.4001705020945628, 0.3930644655688551, 0.392335542092979, 0.392432524331171, 0.3927347253426151] [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1] *****bestHypermeterFunc Completed***** for best_alpha = 0.01 For values of best alpha = 0.01 The train log loss is: 0.3605788274237615 For values of best alpha = 0.01 The cross validation log loss is: 0.392335542092979 For values of best alpha = 0.01 The test log loss is: 0.37951616221945994</p> <table border="1"> <thead> <tr> <th>Metric</th><th>Train_value</th><th>CV_value</th><th>Test_value</th></tr> </thead> <tbody> <tr> <td>Log-Loss</td><td>0.361</td><td>0.392</td><td>0.38</td></tr> <tr> <td>Precision</td><td>0.652</td><td>0.75</td><td>0.5</td></tr> <tr> <td>Recall</td><td>0.097</td><td>0.135</td><td>0.079</td></tr> <tr> <td>F1-Score</td><td>0.168</td><td>0.229</td><td>0.137</td></tr> <tr> <td>Accuracy</td><td>0.846</td><td>0.853</td><td>0.838</td></tr> </tbody> </table> <p>*****</p>	Metric	Train_value	CV_value	Test_value	Log-Loss	0.361	0.392	0.38	Precision	0.652	0.75	0.5	Recall	0.097	0.135	0.079	F1-Score	0.168	0.229	0.137	Accuracy	0.846	0.853	0.838	
Metric	Train_value	CV_value	Test_value																								
Log-Loss	0.361	0.392	0.38																								
Precision	0.652	0.75	0.5																								
Recall	0.097	0.135	0.079																								
F1-Score	0.168	0.229	0.137																								
Accuracy	0.846	0.853	0.838																								

			
2.	Tune the Model Accuracy	<p>Training Accuracy – 80% linear</p> <p>for alpha = 1e-06 Log Loss : 0.6143143363763562</p> <p>for alpha = 1e-05 Log Loss : 0.6143143337374273</p> <p>for alpha = 0.0001 Log Loss : 0.6142136356700172</p> <p>for alpha = 0.001 Log Loss : 0.5871537845084024</p> <p>for alpha = 0.01 Log Loss : 0.5817866567511131</p> <p>for alpha = 0.1 Log Loss : 0.583009140383031</p> <p>for alpha = 1 Log Loss : 0.5832164963645</p>  <p>linear [0.6143143363763562, 0.6143143337374273, 0.6142136356700172, 0.5871537845084024, 0.5817866567511131, 0.583009140383031, 0.5832164963645] [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1] *****bestHypermerFunc Completed*****</p> <p>for best_alpha = 0.01 For values of best alpha = 0.01 The train log loss is: 0.5278361295624395</p> <p>For values of best alpha = 0.01 The cross validation log loss is: 0.5817866567511131</p> <p>For values of best alpha = 0.01 The test log loss is: 0.5474462705866496</p> <p>----- Confusion Matrix -----</p>	 



Metric	Train_value	CV_value	Test_value
Log-Loss	0.528	0.582	0.547
Precision	0.732	0.289	0.329
Recall	0.774	0.622	0.698
F1-Score	0.752	0.394	0.447
Accuracy	0.745	0.692	0.721

Starting Get Important Features Function
linear

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

# Load data
data = pd.read_csv('data.csv')

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(data[['feature1', 'feature2', 'feature3', 'feature4', 'feature5']], data['target'], test_size=0.2, random_state=42)

# Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Train Logistic Regression model
model = LogisticRegression()
model.fit(X_train, y_train)

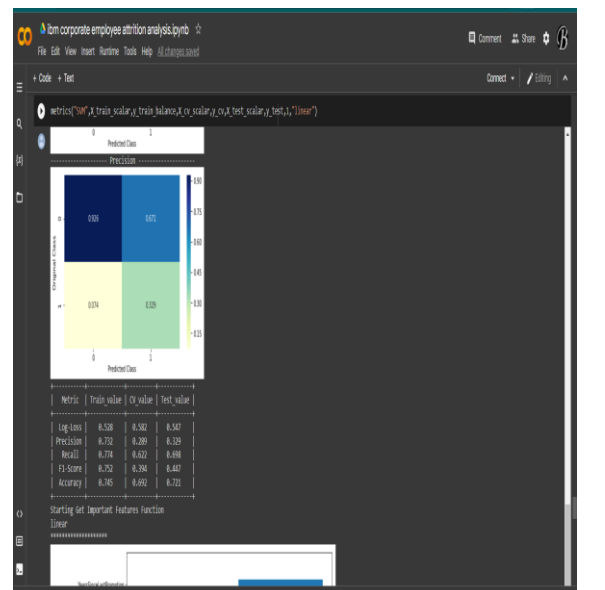
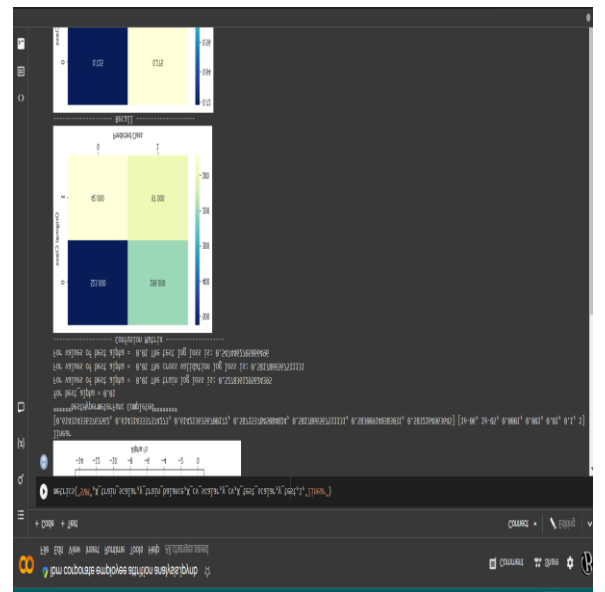
# Predict on test set
y_pred = model.predict(X_test)

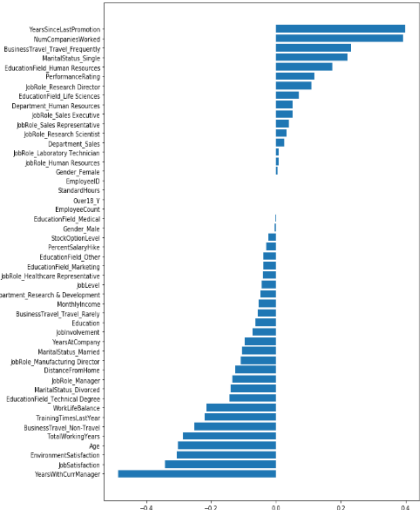
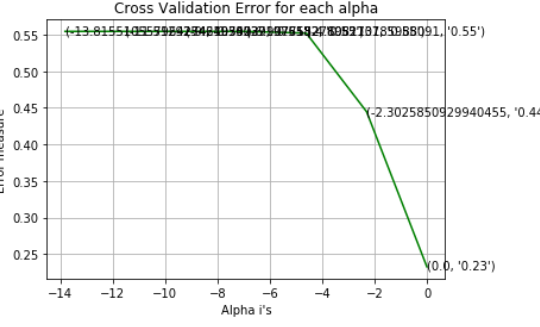
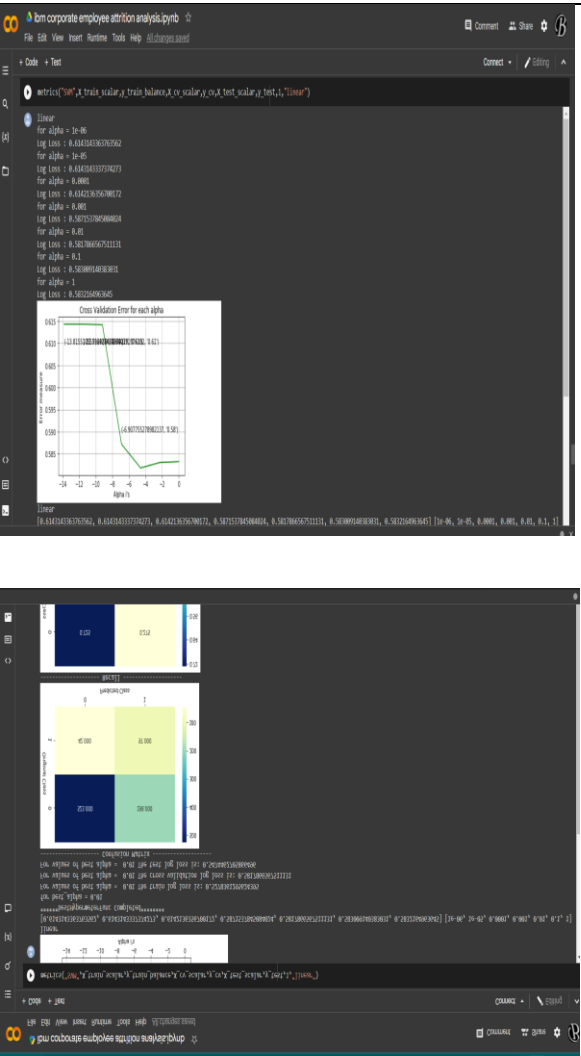
# Evaluate model
cm = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(cm)

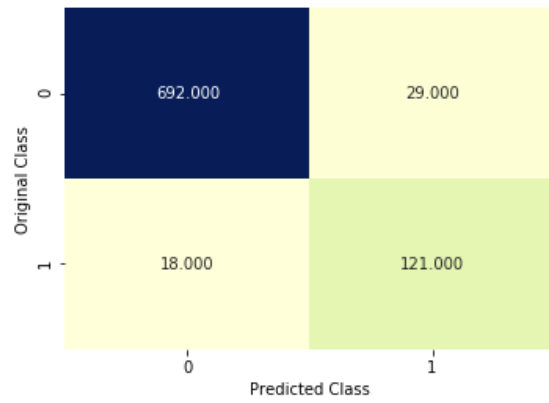
print('Classification Report:')
print(classification_report(y_test, y_pred))

print('Accuracy Score:')
print(accuracy_score(y_test, y_pred))

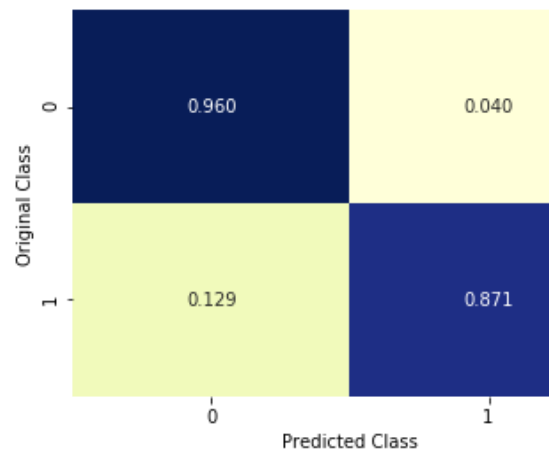
```



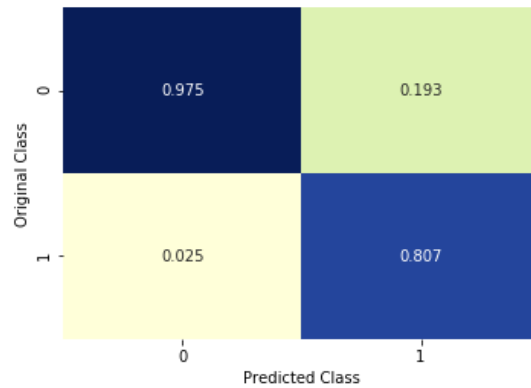
			
3.	Confidence Score (Only Yolo Projects)	<p>Validation Accuracy – 76%</p> <p>Confidence Score – for alpha = 1e-06 Log Loss : 0.5544280465616621 for alpha = 1e-05 Log Loss : 0.5544279653425883 for alpha = 0.0001 Log Loss : 0.5544279626082605 for alpha = 0.001 Log Loss : 0.554427965596575 for alpha = 0.01 Log Loss : 0.5520345548670238 for alpha = 0.1 Log Loss : 0.4446536394141772 for alpha = 1 Log Loss : 0.23267354410235444</p> <p>Cross Validation Error for each alpha</p>  <p>poly [0.5544280465616621, 0.5544279653425883, 0.5544279626082605, 0.554427965596575, 0.5520345548670238, 0.4446536394141772, 0.23267354410235444] [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1] *****bestHypermeterFunc Completed***** for best_alpha = 1 For values of best alpha = 1 The train log loss is: 0.05843126754768401 For values of best alpha = 1 The cross validation log loss is: 0.23267354410235444 For values of best alpha = 1 The test log loss is: 0.19790587725653494 ----- Confusion Matrix -----</p>	



----- Recall -----



----- Precision -----

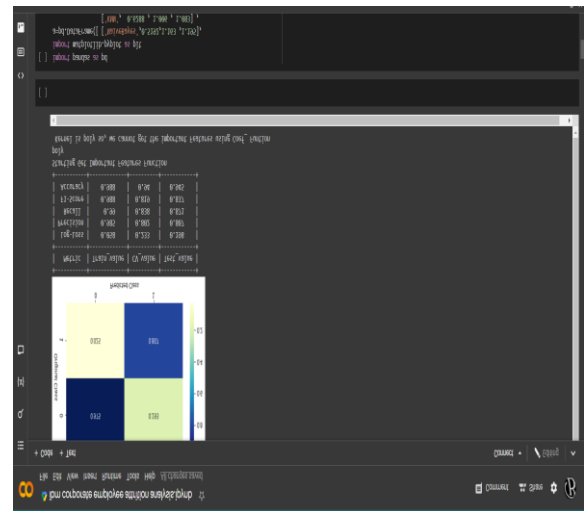
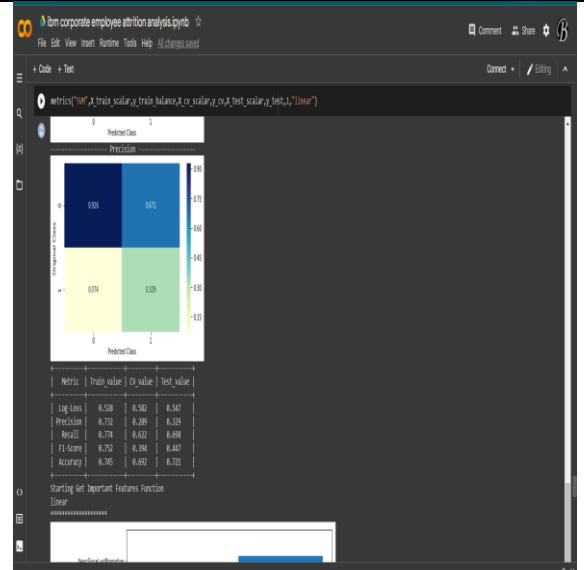


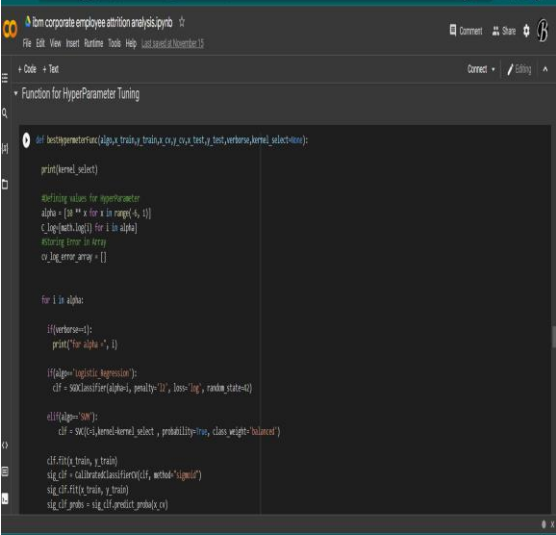
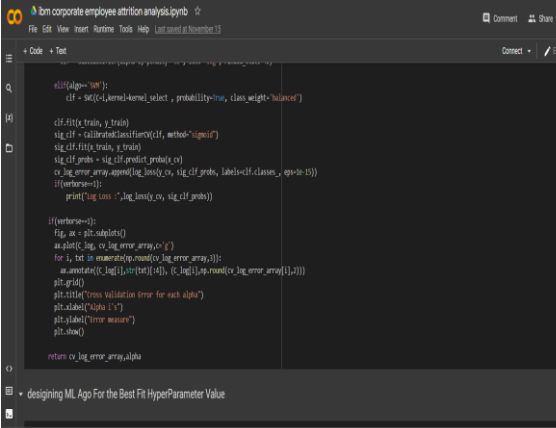
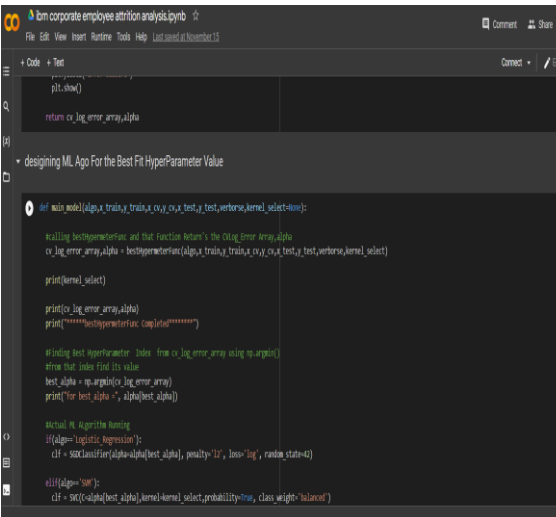
Metric	Train_value	CV_value	Test_value
Log-Loss	0.058	0.233	0.198
Precision	0.985	0.802	0.807
Recall	0.99	0.838	0.871
F1-Score	0.988	0.819	0.837
Accuracy	0.988	0.94	0.945

Starting Get Important Features Function

poly

Kernel is poly so, we cannot get the important Features using Coef_ Funtion



			 <pre> def besthyperparameter(algux_train_x, algux_train_y, algux_test_x, algux_test_y, verbose, kernel_select=None): print(kernel_select) #defining values for hyperparameter alpha = [10**x for x in range(-4, 1)] C_log=[math.log(i) for i in alpha] #storing error in array cv_log_error_array = [] for i in alpha: if(verbose==1): print("For alpha = ", i) if(algux=='logistic regression'): clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42) elif(algux=='svm'): clf = SVC(kernel=kernel_select, probability=True, class_weight='balanced') clf.fit(x_train, y_train) sig_clf = CalibratedClassifierWrapper(clf, method='sigmoid') sig_clf.fit(x_train, y_train) sig_clf_proba = sig_clf.predict_proba(x) if(verbose==1): print("log loss = ", log_loss(y_cv, sig_clf_proba)) if(verbose==0): fig, ax = plt.subplots() ax.plot(C_log, cv_log_error_array, 'r') for i, text in enumerate(np.round(cv_log_error_array, 4)): ax.annotate(text, (C_log[i], 0.015*(1-i)), (C_log[i], log_loss(y_cv, sig_clf_proba)), (i, 1)) plt.grid() plt.title("Cross validation error for each alpha") plt.xlabel("alpha (r^2)") plt.ylabel("error measure") plt.show() return cv_log_error_array, alpha </pre>
			 <pre> elif(algux=='svm'): clf = SVC(kernel=kernel_select, probability=True, class_weight='balanced') clf.fit(x_train, y_train) sig_clf = CalibratedClassifierWrapper(clf, method='sigmoid') sig_clf.fit(x_train, y_train) sig_clf_proba = sig_clf.predict_proba(x) cv_log_error_array.append(log_loss(y_cv, sig_clf_proba, labels=clf.classes_, eps=1e-15)) if(verbose==1): print("log loss = ", log_loss(y_cv, sig_clf_proba)) if(verbose==0): fig, ax = plt.subplots() ax.plot(C_log, cv_log_error_array, 'r') for i, text in enumerate(np.round(cv_log_error_array, 4)): ax.annotate(text, (C_log[i], 0.015*(1-i)), (C_log[i], log_loss(y_cv, sig_clf_proba)), (i, 1)) plt.grid() plt.title("Cross validation error for each alpha") plt.xlabel("alpha (r^2)") plt.ylabel("error measure") plt.show() return cv_log_error_array, alpha </pre> <p>designing ML Algo For the Best Fit HyperParameter Value</p>
			 <pre> def main_model(algux_train_x, algux_train_y, algux_test_x, algux_test_y, verbose, kernel_select=None): #calling besthyperparameter and that function return's the Cross Error Array, alpha cv_log_error_array, alpha = besthyperparameter(algux_train_x, algux_train_y, algux_test_x, algux_test_y, verbose, kernel_select) print(kernel_select) print(cv_log_error_array, alpha) print("*****besthyperparameter Completed*****") #finding best hyperparameter index from cv_log_error_array using np.argmin() #from that index find its value best_alpha = np.argmin(cv_log_error_array) print("For best alpha = ", alpha[best_alpha]) #actual ML Algorithm training if(algux=='logistic regression'): clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42) elif(algux=='svm'): clf = SVC(kernel=kernel_select, probability=True, class_weight='balanced') </pre>

			<div><div>bm corporate employee attrition analysis.ipynb</div><div><div>File Edit View Insert Runtime Tools Help Last saved at November 15</div><div><div>+ Code + Text</div><div>Connect</div><div>Editing</div></div></div><div><pre>clf(alpha=500): clf = SVC(kernel='rbf', gamma=0.001, class_weight='balanced') clf.fit(x_train, y_train) sig_clf = CalibratedClassifierCV(clf, method='sigmoid') sig_clf.fit(x_train, y_train) #Predicting output labels for all train/test data predict_train=sig_clf.predict(x_train) predict_cv=sig_clf.predict(x_test) predict_test=sig_clf.predict(x_test) #Finding Logloss Metrics for all data using predictProb values predictProb_train = sig_clf.predict_proba(x_train) logloss_train=log_loss(y_train, predictProb_train, labels=clf.classes_, eps=1e-15) print("For values of best alpha = ", alpha[best_alpha], "The train log loss is:",logloss_train) predictProb_cv = sig_clf.predict_proba(x_test) logloss_cv=log_loss(y_test, predictProb_cv, labels=clf.classes_, eps=1e-15) print("For values of best alpha = ", alpha[best_alpha], "The cross validation log loss is:",logloss_cv) predictProb_test = sig_clf.predict_proba(x_test) logloss_test=log_loss(y_test, predictProb_test, labels=clf.classes_, eps=1e-15) print("For values of best alpha = ", alpha[best_alpha], "The test log loss is:",logloss_test) if(verbose==1): #calling plotting function plot_confusion_matrix(y_test,predict_test)</pre></div></div>
--	--	--	--